

Astar training 2025

This a notebook for the Astar training 2025. You are guided to train a simple classifier on the cifar-10 dataset.

Instruction

*Please try to read through this notebook, and fill in the blank according to the instructions given after **TODO**.*

You may download the notebook or excute directly on kaggle.

Submission

submit .pdf converted from this notebook after completion

– Fan Zhenyi

code Requirements

This notebook requires the following packages

excute the following command to install the required packages

```
1 !pip install torch torchvision
2 !pip install matplotlib
3 !pip install numpy
4 !pip install tqdm
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (2.4.1+cu121)
Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages (0.19.1+cu121)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch) (3.16.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch) (4.12.2)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch) (1.13.3)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch) (3.3)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch) (3.1.4)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch) (2024.6.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torchvision) (1.26.4)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.10/dist-packages (from torchvision) (10.4.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch) (2.1.5)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy->torch) (1.3.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.54.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.7)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.26.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib) (1.16.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.26.4)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.66.5)
```

```
1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 from torch.utils.data import DataLoader, Dataset
5 from torchvision import datasets, transforms
6
7
8 import matplotlib.pyplot as plt
9
10 import numpy as np
11 import pickle
12
13 import os
14
15 from tqdm import tqdm
```

Data preparation

The dataloader is already implemented for you. You can use it to load the cifar-10 dataset.

```

1 class CIFAR10(Dataset):
2     #__init__ is a special method in python classes, it is called when an object is created
3     def __init__(self, root, train=True, transform=None):
4         #root path to the dataset, you may ignore this
5         self.root = root
6
7         #this is a boolean value to indicate whether we are loading the training or test set
8         self.train = train
9
10        #this is the transformation that will be applied to the images
11        #it's none by default, you are required to pass a transform later
12        self.transform = transform
13
14        #checking if the dataset exists, if not download it
15        if not self._check_exists():
16            print("Downloading cifar10 dataset")
17            self.download()
18
19        #load the data
20        #We are going to load the data in memory
21        #Store the images and labels in the self.data and self.targets variables
22        if self.train:
23            self.data, self.targets = self._load_training_data()
24        else:
25            self.data, self.targets = self._load_test_data()
26
27        #this method returns the length of the dataset
28        def __len__(self):
29            return len(self.data)
30
31        #this method returns a sample from the dataset at the given index
32        #this is very important because it allows us to iterate over the dataset
33        def __getitem__(self, index):
34            img, target = self.data[index], int(self.targets[index])
35
36            img = torch.from_numpy(img).float().permute(2,0,1) / 255.0
37
38            if self.transform:
39                img = self.transform(img)
40
41            return img, target
42
43
44        def _check_exists(self):
45            return os.path.exists(os.path.join(self.root, "cifar-10-batches-py"))
46
47        def download(self):
48            if self._check_exists():
49                print("Dataset already exists !!!")
50                return
51            return datasets.CIFAR10(self.root, train=self.train, download=True)
52
53        #this function looks complicated but it's just reading the data from the files
54        def _load_batch(self, file_path):
55            with open(file_path, 'rb') as f:
56                data = pickle.load(f, encoding='bytes')
57            return data[b'data'], data[b'labels']
58
59        def _load_training_data(self):
60            data = []
61            targets = []
62            for i in range(1, 6):
63                file_path = os.path.join(self.root, "cifar-10-batches-py", f"data_batch_{i}")
64                batch_data, batch_labels = self._load_batch(file_path)
65                data.append(batch_data)
66                targets.extend(batch_labels)
67
68            data = np.vstack(data).reshape(-1, 3, 32, 32).transpose(0, 2, 3, 1)
69            return data, np.array(targets)
70
71        def _load_test_data(self):
72            file_path = os.path.join(self.root, "cifar-10-batches-py", "test_batch")
73            data, labels = self._load_batch(file_path)
74            return data.reshape(-1, 3, 32, 32).transpose(0, 2, 3, 1), np.array(labels)
75

```

```

1 #We have define our datasetclass, now we are going to instantiate it
2 #The instantiation will download the dataset and load it in memory, it takes about 1-2 mins
3 train_dataset = CIFAR10(root="data", train=True)
4 test_dataset = CIFAR10(root="data", train=False)
5
6 #simple demonstration __getitem__ method
7 img, target = train_dataset[0] #get the first image in the dataset
8 plt.figure(figsize=(1,1))
9 plt.imshow(img.permute(1,2,0))
10 #TODO: what does permute do?, and why do we need it here?
11 #answer:
12 # PyTorch uses CHW format images while plt takes in HWC format images
13 # Use permute to interchange the dimensions of a tensor to make it compatible with plt
14 #end of you answer
15 plt.title("Original Image")
16 plt.axis('off')
17 print(img.shape, "label: ",target)
18 #the image is a torch tensor (3, 28, 28) and the target is the label of the image
19
20
21 #TODO: define reasonable transformations for the images, you can use the transforms module from torchvision
22 transform = transforms.Compose([
23     transforms.ToPILImage(),
24
25     #TODO: Implement the transformations here
26     transforms.Resize(size=(32, 32)), # Resize the image
27     transforms.RandomRotation(15),
28     transforms.RandomHorizontalFlip(0.5), #Image augmentation
29     #end of your implementation;
30
31     transforms.ToTensor()
32 ])
33
34 train_dataset.transform = transform
35
36 img, target = train_dataset[0] #get the first image in the dataset
37 plt.figure(figsize=(1,1))
38 plt.imshow(img.permute(1,2,0))
39 plt.axis('off')
40 plt.title("Transformed Image")
41 print(img.shape, "label: ",target)
42

```



Downloading cifar10 dataset

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to data/cifar-10-pyt


100%|██████████| 170498071/170498071 [00:03<00:00, 48500386.66it/s]

Extracting data/cifar-10-python.tar.gz to data

torch.Size([3, 32, 32]) label: 6

torch.Size([3, 32, 32]) label: 6

Original Image



Transformed Image



```

1 #apart from test set, we are going to use the training set to create a validation set
2 #we are going to split the training set into two parts
3 train_size = int(0.9 * len(train_dataset))
4 val_size = len(train_dataset) - train_size
5 train_dataset, val_dataset = torch.utils.data.random_split(train_dataset, [train_size, val_size])
6
7 #we are going to use the DataLoader class to create an iterator for our dataset
8 #this iterator will be used to iterate over the dataset in batches
9 #tentatively we are going to use a batch size of 32
10 #TODO: change different batch sizes and see how it affects the training process
11 # Answer: larger batch size, faster training speed, but less times of optimization, i.e. less opportunities to learn from the data.
12 # Why batch? Dealing with data one by one is quite slow, but considering many at a time reduces learning opportunities.
13 train_loader = DataLoader(train_dataset, batch_size=256, shuffle=True, num_workers= 2, pin_memory=True)
14 val_loader = DataLoader(val_dataset, batch_size=256, shuffle=False, num_workers= 2, pin_memory = True)
15 test_loader = DataLoader(test_dataset, batch_size=256, shuffle=False, num_workers= 2, pin_memory = True)
16 # Modify number of workers to suit the hardware.

```

Model

This defines the model architecture. You can use the model as it is or modify it as per your requirements.

```

1 class LinearModel(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.name = "Linear"
5         self.num_inputs = 3*32*32
6         hidden_size = 512
7         num_classes = 10
8         super().__init__()
9         self.linear = nn.Sequential(
10             nn.Linear(self.num_inputs, hidden_size),    # batch_size x 784 -> batch_size x 512
11             nn.ReLU(), #activation function              # batch_size x 512 -> batch_size x 512
12             nn.Linear(hidden_size, num_classes)         # batch_size x 512 -> batch_size x 10
13         ) #nn.Sequential is a container for other layers, it applies the layers in sequence
14
15     #forward is the method that defines the forward pass of the network
16     #not rigorously: model.forward(x) = model(x)
17     def forward(self, x):
18         x = x.view(-1, self.num_inputs) # flatten the image from 3x32x32 to 3072
19         x = self.linear(x)
20         return x
21
22 class CNNModel(nn.Module):
23     def __init__(self):
24         super().__init__()
25         self.name = "CNN"
26         self.conv = nn.Sequential(
27             #TODO: write the size of the input and output of each layer e.g
28             nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1), #input: 3x32x32, output: 32x32x32
29             nn.ReLU(),
30             nn.MaxPool2d(kernel_size=2, stride=2),
31             nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
32             nn.ReLU(),
33             nn.MaxPool2d(kernel_size=2, stride=2),
34             nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
35             nn.ReLU(),
36             nn.MaxPool2d(kernel_size=2, stride=2)
37         )
38         self.fc = nn.Sequential(
39             nn.Linear(128*4*4, 512),
40             nn.ReLU(),
41             nn.Linear(512, 10)
42         )
43
44     def forward(self, x):
45         x = self.conv(x)
46         x = x.view(-1, 128*4*4)
47         x = self.fc(x)
48         return x
49
50     def getFeature(self, x):
51         x = self.conv(x)
52         feat = x.view(-1, 128*4*4) #this is the 128*4*4 feature
53         return feat
54

```

```

55
56 class YourModel(nn.Module):
57     def __init__(self, input_shape: int = 3, hidden_units: int=20):
58         super().__init__()
59         self.name = "Your_model"
60         #TODO: define your model here
61         self.input_size = input_shape
62         self.hidden_units = hidden_units
63         self.conv_block = nn.Sequential(
64             nn.Conv2d(in_channels = input_shape, out_channels=hidden_units, kernel_size=3, stride=1, padding=1),
65             nn.ReLU(),
66             nn.Conv2d(in_channels = hidden_units, out_channels=hidden_units, kernel_size=3, stride=1, padding=1),
67             nn.ReLU(),
68             nn.MaxPool2d(kernel_size = 2),
69             nn.Conv2d(in_channels = hidden_units, out_channels=hidden_units, kernel_size=3, stride=1, padding=1),
70             nn.ReLU(),
71             nn.Conv2d(in_channels = hidden_units, out_channels=hidden_units, kernel_size=3, stride=1, padding=1),
72             nn.ReLU(),
73             nn.MaxPool2d(kernel_size = 2)
74         )
75         self.classifier = nn.Sequential(
76             nn.Flatten(),
77             nn.Linear(in_features= 1280, # To be modified by error message
78                       out_features = 10)
79         )
80
81     def forward(self, x):
82         #TODO: implement the forward pass of your model
83         return self.classifier(self.conv_block(x))

```

Training

This is the training loop. You can modify the training loop as per your requirements.

The training takes some time. You can do other tasks while the training is in progress.

you may use the gpu on kaggle or colab to speed up the training process. <https://www.kaggle.com/discussions/general/97939>

```

1 def training(model: torch.nn.Module,
2             train_loader: torch.utils.data.DataLoader,
3             val_loader: torch.utils.data.DataLoader,
4             epochs: int = 20,
5             least_early_stopping: int = 10,
6             device = "cuda"):
7     #defining the loss function
8     criterion = nn.CrossEntropyLoss()
9
10    #defining the optimizer
11    #TODO: try to modify the optimizer and see how it affects the training process
12    # Not all optimizer suits for classification task ...
13    optimizer = optim.Adam(model.parameters(), lr=0.01)
14
15    best_accuracy = 0
16    #early stopping
17    early_stopping = least_early_stopping
18    early_stopping_counter = 0
19
20
21    #TODO: adjust the number of epochs and see how it affects the training process
22    epochs = 20
23    for epoch in range(epochs):
24        #training
25        for images, labels in tqdm(train_loader):
26            images = images.to(device)
27            labels = labels.to(device)
28            #forward pass
29            outputs = model(images)
30            #calculate the loss
31            loss = criterion(outputs, labels)
32            #zero the gradients
33            optimizer.zero_grad() #ensure that the gradients are zero
34            #backward pass
35            loss.backward()
36            #optimize
37            optimizer.step()
38        #validation
39        total = 0
40        correct = 0
41        for images, labels in val_loader:
42            images = images.to(device)
43            labels = labels.to(device)
44            outputs = model(images)
45            _, predicted = torch.max(outputs, 1)
46            total += labels.size(0)
47            correct += (predicted == labels).sum().item()
48        accuracy = correct / total
49
50        #TODO: implement early stopping
51        #what is early stopping? https://en.wikipedia.org/wiki/Early_stopping
52        if accuracy < early_stopping and epoch > early_stopping:
53            break
54        #end of early stopping
55
56        if accuracy > best_accuracy:
57            best_accuracy = accuracy
58            #save the model
59            torch.save(model.state_dict(), f"best_model_{model.name}.pth")
60            print(f"Epoch: {epoch}, Loss: {loss.item()}, Accuracy: {accuracy} ***")
61        else:
62            print(f"Epoch: {epoch}, Loss: {loss.item()}, Accuracy: {accuracy}")
63

```

```

1 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
2 print("using device:", device)
3 #instantiating the model
4 #TODO: change the model to the CNNModel and Your model, and evaluate the performance.
5 model = CNNModel()
6 model = model.to(device)
7 print(f"training {model.name} model")
8 training(model, train_loader, val_loader, 25, 10)
9
10 print("-----")
11
12 yourmodel = YourModel()
13 yourmodel.to(device)
14 print(f"training {yourmodel.name} model")
15 training(yourmodel, train_loader, val_loader, 25, 10)

```

```

→ using device: cuda
training CNN model
100%|██████████| 159/159 [00:16<00:00, 9.68it/s]
Epoch: 0, Loss: 1.6888031959533691, Accuracy: 0.3684444444444446 ***
100%|██████████| 159/159 [00:15<00:00, 10.57it/s]
Epoch: 1, Loss: 1.6128902435302734, Accuracy: 0.4586666666666667 ***
100%|██████████| 159/159 [00:15<00:00, 10.60it/s]
Epoch: 2, Loss: 1.3689610958099365, Accuracy: 0.4831111111111111 ***
100%|██████████| 159/159 [00:14<00:00, 10.61it/s]
Epoch: 3, Loss: 1.4349294900894165, Accuracy: 0.4973333333333335 ***
100%|██████████| 159/159 [00:15<00:00, 10.41it/s]
Epoch: 4, Loss: 1.4723223447799683, Accuracy: 0.5206666666666667 ***
100%|██████████| 159/159 [00:14<00:00, 10.65it/s]
Epoch: 5, Loss: 1.4026802778244019, Accuracy: 0.5493333333333333 ***
100%|██████████| 159/159 [00:15<00:00, 10.56it/s]
Epoch: 6, Loss: 1.2192330360412598, Accuracy: 0.546
100%|██████████| 159/159 [00:15<00:00, 10.59it/s]
Epoch: 7, Loss: 1.2267341613769531, Accuracy: 0.5377777777777778
100%|██████████| 159/159 [00:16<00:00, 9.43it/s]
Epoch: 8, Loss: 1.1706657409667969, Accuracy: 0.5653333333333334 ***
100%|██████████| 159/159 [00:15<00:00, 10.51it/s]
Epoch: 9, Loss: 1.1283478736877441, Accuracy: 0.5317777777777778
100%|██████████| 159/159 [00:14<00:00, 10.66it/s]
Epoch: 10, Loss: 1.2462068796157837, Accuracy: 0.5057777777777778
100%|██████████| 159/159 [00:15<00:00, 10.57it/s]
-----
training Your_model model
100%|██████████| 159/159 [00:15<00:00, 10.00it/s]
Epoch: 0, Loss: 1.836076021194458, Accuracy: 0.3102222222222223 ***
100%|██████████| 159/159 [00:15<00:00, 10.47it/s]
Epoch: 1, Loss: 1.7102361917495728, Accuracy: 0.378 ***
100%|██████████| 159/159 [00:16<00:00, 9.63it/s]
Epoch: 2, Loss: 1.5984342098236084, Accuracy: 0.414 ***
100%|██████████| 159/159 [00:15<00:00, 10.40it/s]
Epoch: 3, Loss: 1.296134114265442, Accuracy: 0.4646666666666667 ***

```